

Training a neural network.

In the last section we looked at how neural networks are able to perform difficult inference tasks: through the careful choice of weights and biases they can produce complicated non-linear functions. This does not explain their power, there are lots of ways to produce complicated function, the challenge is learning the functions from data and we haven't discussed how neural networks are trained. We will discuss that here, but in a way that won't explain the full power of these networks; there is an unreasonable effectiveness to learning in deep-learning neural networks, deep here means networks with many layers and the role deepness plays in this unreasonable effectiveness is one of the mysteries.

Sometimes too much is made of the fact that we don't fully understand how deep learning networks work, or, rather, why they work so well. However, one of the glories of engineering is that something can be useful even if you don't fully understand it and a mystery, of course, is exciting for scientists, it is a door you can hope to pass through, a door into the dark.

Anyway, we have already touched on the mechanism used to train a neural network. Basically, if you have an objective function, a measure how well something is working and some parameters that adjust the performance of the network as quantified using the objective function, then you can use some hill-climbing or optimisation algorithm to improve the performance and, in the case of neural networks, that algorithm is a variant of stochastic gradient descent¹. We have already seen the essential ingredients for this when we looked at logistic regression; basically we use cross-entropy loss as the objective and optimize it.

¹Here, as in a number of places in mathematics, we are very sloppy about which way we think of as up! We talk about hill-climbing algorithms and stochastic gradient descent, even though descending means going down, not climbing. This happens because, of course, the notion of up and down is metaphorical, are we minimising an error function or maximising some sort of likelihood function? Usually, by convention these days we head downwards, so if we are using a likelihood function we add a minus, but it is just a convention and all our terminology hasn't quite caught up. The other place you'll often see up versus down confusion is when talking about search trees where you usually imagine the tree with the roots at the top!

Cross-entropy loss and SGD

Lets consider a typical situation, you have input \mathbf{x} and output \mathbf{y} so that $\mathbf{y}(\mathbf{x}; \theta)$ and the θ are a set of parameters. In the most straight-forward case these parameters will be weights and biases for a neural network. Typically there will be a dataset

$$\mathcal{D} = \{(\mathbf{x}^a, \mathbf{y}^a) | a = 1 \dots N\} \quad (1)$$

so, here, there are N data points and I am using a superscript as the data index, a subscript will be used for the components of the vector², so the first data point in the list is

$$\mathbf{x}^1 = (x_1^1, x_2^1, \dots, x_{n_x}^1) \quad (2)$$

and n_x here is the size of the input, we will use n_y for the size of the output. For a classification problem the \mathbf{y}^a will be *one hot vectors*³, which means they will have a single one value corresponding to the class of the item and zeros for all the other entries, so if the first item in the data set belongs to class k then

$$\mathbf{y}^1 = (0, 0, \dots, 1, \dots, 0, 0) \quad (3)$$

Now we are using notation that is potentially confusing, \mathbf{y} without a superscript is a function, so $\mathbf{y}(\mathbf{x}^a)$ is a set of numbers calculated using the neural network based on the input \mathbf{x}^a .

We are ready to define the cross entropy loss, it is

$$\mathcal{L} = -\frac{1}{N} \sum_{a=1}^N \sum_{i=1}^{n_y} y_i^a \log p_i^a \quad (4)$$

²I have previously owned up to the inconsistency of how I use indices; computer scientists reasonably don't use superscripts so much since in computer code indices are indices!

³Almost all of machine learning for classification uses one hot vectors; at first it seems profligate, allowing denser coding in the outside later would result in much shorter output vectors, indeed logarithmically shorter, but this doesn't work, the corresponding outputs are just much harder to learn and the idea of considering the output as giving probabilities to compare to the one-hot target output would be very complicated for a dense representation. In practice, the neural network learns a dense representation itself for the layer immediately before the output layer and a matrix of weights translating this the typically very large output vector.

where

$$p_i^a = \sigma_i(\mathbf{y}^a) \quad (5)$$

is the softmax. The second sum in the loss function has only one term since y_i^a is mostly zero:

$$-\sum_{i=1}^{n_y} y_i^a \log p_i^a = -\log p_k^a \quad (6)$$

if the a th item belongs to class k . We can see that minimizing this is good, we want this probability to be close to one and $\log 1 = 0$ while the log of numbers less than one are negative.

The standard way to minimize this is some version of gradient flow. We discussed gradient flow before when looking at regression; the idea is that we change all the parameters by a small amount along a direction that points down the L landscape:

$$\theta \leftarrow \theta - \eta \frac{\partial \mathcal{L}}{\partial \theta} \quad (7)$$

where θ is a parameter of the model, one of the Ws or bs , η is the learning rate, a small number which allows the algorithm to move down towards the minimum, hopefully fast enough not to take forever, but with steps small enough not to overshoot and cause unstable, oscillating behaviour. Now the problem is how to calculate all the derivatives.

At first this looks like it would be impossibly difficult, the models are very complicated and so calculating all the derivatives would be a lot of work. It is a lot of work, but we do not have to do it, the computer does it for us. This is one of the unsung miracles of our current age, basically, using the chain rule we know how to differentiate functions that can be written as one thing after another: if $y(x) = f(g(x))$, so you get u from x first by doing g to it and then f to the result, we know that the rate of change of y with x is the multiple of how much f changes with g by how much g changes with x :

$$\frac{dy}{dx} = \frac{df}{dg} \frac{dg}{dx} \quad (8)$$

Now, when we programme the neural network we are essentially telling the computer how do the calculation step-by-step, ultimately to calculate the computer breaks these instructions into still smaller steps, a series of elementary operations. The programme can keep track of these elementary steps and use the chain rule, and the product rule, to work out the derivative, this is called **autograd** and is an essential component of what has made

modern machine learning possible, until recently, the human did the calculus, now the machine does it, allowing us to use much more complicated networks and to change them without having to redo all the mathematics.

These is another problem; working out \mathcal{L} as defined would take a long time, for neural networks to work we need to have a lot of data and so the calculation of the likelihood would take a huge number of calculations for just one small gradient step. However, it turns out that using a much smaller subset will give some reasonable sense of which direction to move in; approximating \mathcal{L} from a *batch* of maybe ten or a 100 data points is often enough, the key thing is to use a different batch for each gradient step until you have gone through all the data, an achievement often called an *epoch*. Hence, we calculate

$$\mathcal{L}(B) = -\frac{1}{N_B} \sum_{a \in B} \sum_{i=1}^{n_y} y_i^a \log p_i^a \quad (9)$$

where B is a batch, some N_B sized subset of the total data and then do

$$\theta \leftarrow \theta - \eta \frac{\partial \mathcal{L}(B)}{\partial \theta} \quad (10)$$

for all the parameters θ before going on to the next batch.

Weirdly, this process of using batches does not just make optimization quicker, it also make better. The \mathcal{L} landscape is not a simple bowl; unlike simpler cases like regression where we expect an easy optimisation processes, the neural network is designed to produce non-linear maps and so, inevitably, it has a loss function that may have local minima and other annoying features: this is sketched out in Fig. 1. Now if batches are used instead each gradient step will correspond to a different $\mathcal{L}(B)$ each approximating the objective, this is illustrated in Fig. 2. This introduces noise which should affect some features more than others, meaning that the descent is less likely to get stuck, the use of batches adds a little randomness which actually seems to make the optimization more robust; as a side-note the idea that SGD helps you avoid getting stuck is probably a small bit misleading, the gradient flow proceeds in a very high-dimensional space and so local minima are probably rare compared to narrow and tortuous valleys and the stochastic element is probably most useful in avoiding those; so the cartoon in the figure is like a slice through the high-dimensional space and the point holds even if the directions are at minima.

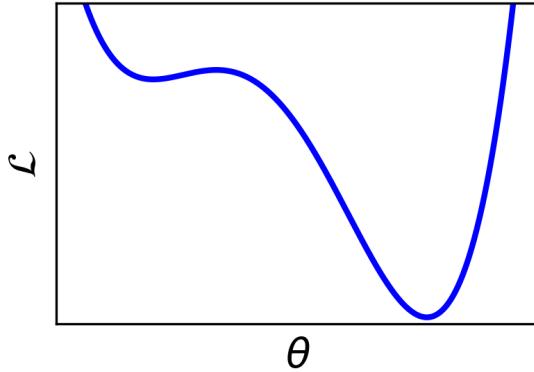


Figure 1: A cartoon of an objective function, the x -axis stands in for all the many parameter directions. There is a local minimum on the left, if the system started off to the left of that and used gradient descent it might get stuck in that local minimum. This local minimum is intended as a stand-in for the hazards, like thin valleys and saddle-points with only a few downward directions that might slow down or ruin gradient descent.

With this element of randomness this approach is called *stochastic gradient descent*, often referred to by its acronym SGD and is a very standard training algorithm. There are variations we won't explore further here that try to do things like improve the size of the gradient step to take into account the local behaviour of \mathcal{L} ; the most common of these approaches is probably *Adam*. Beyond that there are a huge number of techniques for improving training; for example, in *dropout* some of the nodes picked at random are switched off for each training step, this appears to make training robust and improve behaviour on data points not included in the training set.

This is not at all the end of the story. One obvious problem is that you might expect the gradients to be very low when you are far from the true solution; two different very wrong answers will produce more or less the same error, whereas the gradient will be high when we are close to the minimum. This is exactly what we don't want, it means we move slowly when we are far away and precision isn't required and quickly when we are close and it is. This implies some sort of dynamical adjustment should be made to the learning rate and, indeed, there are many approaches to that, including the very popular ADAM optimiser; though, strictly, ADAM does more than

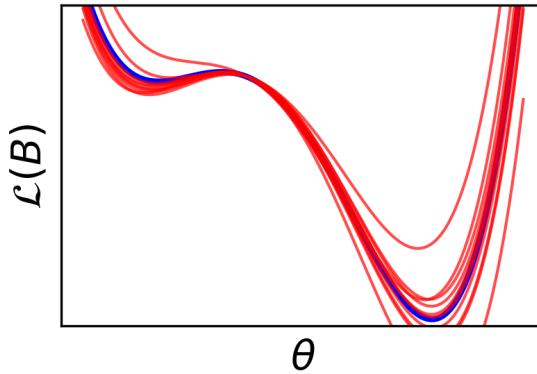


Figure 2: A cartoon of an objective function with batches. Here the blue line represents \mathcal{L} and the red lines $\mathcal{L}(B)$ for different B . The local minimum isn't always a minimum and moves around a lot.

that by including contributions from the second derivative. A lot of current effort in machine learning is focused on issues of how to do learning in an efficient, stable and useful way: adjusting learning steps and rescaling weights so that some don't 'take over' by become too large, for example.